Lab 08

Basic Image Processing Algorithms
Fall 2024

K-means Algorithm

⊙ *K-means* algorithm (MacQueen'67): a heuristic method

- Each cluster is represented by the centre of the cluster and the algorithm converges to stable centroids of clusters.

- K-means algorithm is the simplest partitioning method for clustering analysis and widely used in data mining applications.

- Each sample will belong to the cluster with the nearest mean.

⊙ The objective is to minimize the within-cluster sum of squares:

$$\operatorname*{argmin}_{S} \sum_{i=1}^{K} \sum_{x \in S_i} d^2(x, \mu_i), \qquad d^2(x, \mu_i) = \|x - \mu_i\|^2 = \sum_{n=1}^{N} (x_n - \mu_{in})^2$$

- where $x \in \mathbb{R}^N$ are the data samples, μ_i is the mean (prototype) of the points in the cluster S_i $(i = 1 \dots K)$.

K-means Algorithm

⊙ Given the cluster number K, the *K-means* algorithm is carried out in three steps after initialization:

1) **Initialisation**: set the K cluster seed points (randomly)

2) **Assignment step**: Assign each object to the cluster of the nearest seed point measured with a specific distance metric

3) **Update step**: Compute new seed points as the centroids of the clusters of the current partition (the centroid is the centre, i.e., *mean point*, of the cluster)

$$\mu_i^{(t+1)} = \frac{1}{\left| S_i^{(t)} \right|} \sum_{x_j \in S_i^{(t)}} x_j$$

4) **Go back** to Step 1), stop when no more new assignment (i.e., membership in each cluster no longer changes)

Spaces & dimensions (in pure math)

It is important to understand all the spaces and their dimensionalities in this task.

Consider an S space: $S \subset \mathbb{R}^n$

Every element of S is a vector with n coordinates:

$$\mathbf{x}_a = [x_a^1, x_a^2, ..., x_a^n] \in S$$

The S_i subsets of S are the clusters: $S_i \subset S$

$$\bigcup_{i=1}^{k} S_i = S \quad \text{and } S_i \cap S_j = \emptyset \quad \forall\, i \neq j$$

and where $i = 1, 2, ..., \boxed{k}$

These are the 'k'-s in the term 'k-means'. They tell you the number of clusters.

Also, there are μ vectors representing the mean values aka the centroids of every cluster:

$$\mu_i = [\mu_i^1, \mu_i^2, ..., \mu_i^n] \in S \quad\quad i = 1, 2, ..., \boxed{k}$$

Spaces & dimensions (in MATLAB)

Let us translate the terms of the previous slide into MATLAB.

Consider an S <u>space</u>, represented by a *matrix*.

Every *row* of S is a <u>vector</u> with n items:

  `S(1,:) = x_1 = [x_11 x_12 … x_1n]`

The S_i subsets of S are the <u>clusters</u>. Their rep-representations are stored in a look-up-table (LUT). The index represents the index, the value represents the cluster # of a row vector x_a of S.

Also, the μ mean vectors are stored in a matrix similar to S, denoted by M. Its elements are

  `M(j,:) = mu_j = [mu_j1 … mu_jn]`

S =
← x_1
← x_2
← x_3
⋮
m
n

LUT =
m

M =
← mu_1
← mu_2
⋮
← mu_k
k
n

This is the 'k' we praise :)

5

So what is stored in the LUT?

The LUT does the vector-cluster mapping.

The LUT is a vector. It has as many elements as the number of vectors in the space S.

Every element of the LUT has an *index* and a *value.*

The value at position j tells us which cluster does vector $\mathbf{x}_j$ belong to.

Index: 1   2   3   4   5

LUT = | 1 | 3 | 2 | 1 | 3 |

$m = 5$

6

Now please

download the 'Lab 08' code package

from the

[moodle system](#)

Exercise 1

Implement the function `step1_initialization` **in which:**

- The function has 2 inputs and 2 outputs:

  Inputs:

  - S   set of points to be clustered
  - k   number of clusters

  Outputs:

  - LUT   the assignment vector
  - M   the matrix of centroids

The function should initialize LUT and M as described on the next slide!

After implementation, test your function with the **script** `test1_initialization`!

Step 1: Initialization

S =

- initialize the `LUT` (as an `1 × m` vector, filled with zeros),
- initialize the `M` matrix (as a `k × n` matrix, filled with zeros),
- choose `k`-many vectors as the initial cluster center points and store them in `M`.

$\longrightarrow \mu_1$

$\longrightarrow \mu_2$

$\longrightarrow \mu_3$

Now (to be able to reproduce the results)
we are NOT using random initialization
but an equidistant distribution!

$\longrightarrow \mu_4$

$\lfloor \cdot \rfloor$ means `floor()`

We choose every $\left\lfloor \dfrac{m}{k} \right\rfloor$ -th element as an initial center.

$\longrightarrow \mu_5$

E.g: $m = 16$
$k = 5$ $\Rightarrow$ $\left\lfloor \dfrac{16}{5} \right\rfloor = 3$

The first and then every 3rd element of S will be selected to be an element of `M`.

9

Exercise 2

Implement the function `step2_assignment` **in which:**

● The function has 4 inputs and 1 output:

Inputs:

- ○ S   set of points to be clustered
- ○ k   number of clusters
- ○ LUT   the assignment vector
- ○ M   the matrix of centroids

Outputs:

- ○ LUT   the **updated** assignment vector

The function should update LUT as described on the next slide!

After implementation, test your function with the **script** `test2_assignment`!

Step 2: Assignment

In this step:

For every `x_i` vector in `S` (i=1..m)
    For every `mu_j` vector in `M` (j=1..k)
        Calculate the distance between `x_i` and `mu_j`:

$$d_{ij} = d^2(\mathbf{x}_i, \mu_j) = \|\mathbf{x}_i - \mu_j\|^2 = \sum_{p=1}^{n}(x_i^p - \mu_j^p)^2$$

    From the calculated `d_ij` distances choose the smallest one,
    and store the index of the minimum in the `LUT` at position `i` :

$$\mathrm{LUT}_i = \arg\min_{j=1..k} d_{ij}$$

Exercise 3

Implement the function `step3_update` in which:

- The function has 4 inputs and 1 output:

    Inputs:

    - S — set of points to be clustered
    - k — number of clusters
    - LUT — the assignment vector
    - M — the matrix of centroids

    Outputs:

    - M — the **updated** matrix of centroids

The function should update M as described on the next slide!

After implementation, test your function with the **script `test3_update`**!

Step 3: Update

In this step:

For every **mu_j** vector in **M** (j=1..k)
  Select every **x** vector of **S** that is assigned to the **j**-th cluster:

  MATLAB hint: You can index a vector logically! If the LUT is a vector and you write **LUT == 1** then this expression will return a logical vector: 1 if the element == 1, 0 otherwise.

  If **A = [1 2 3 1 1 2 1]** then **A == 1** returns **[1 0 0 1 1 0 1]**

  The other trick is that if you index a vector or matrix with a logical vector, the result will be the set of those elements that has the same indices where the logical vector contained 1-s.

  If **B = [1 2 3 4 5 6 7**   then **B(:,[1 0 0 1 1 0 1])** returns **[1 4 5 7** .
      2 0 2 0 1 0 7]   **2 0 1 7]**

  Update **mu_j**: the new value is the mean of the vectors of this cluster:

$$\mu_j^{(t+1)} = \frac{1}{|S_j|} \sum_{\mathbf{x}_i \in S_j} \mathbf{x}_i$$

Exercise 4

Implement the function `mykmeans` in which:

- The function has 2 inputs and 3 outputs:

  Inputs:
  - `S`   set of points to be clustered
  - `k`   number of clusters

  Outputs:
  - `LUT` the **final** assignment vector
  - `M`   the **final** matrix of centroids
  - `iter` the number of iterations in which the convergence was reached

The function should realize the iterative procedure described on the next slide. Please print the number of iterations after the execution of the iterative procedure.

After implementation, test your function with the **script** `test4_mykmeans`!

Pseudo-code of the k-means algorithm

function mykmeans(S, k)

    Initialization step

    while not *converged* and *number of iterations is less than 100*
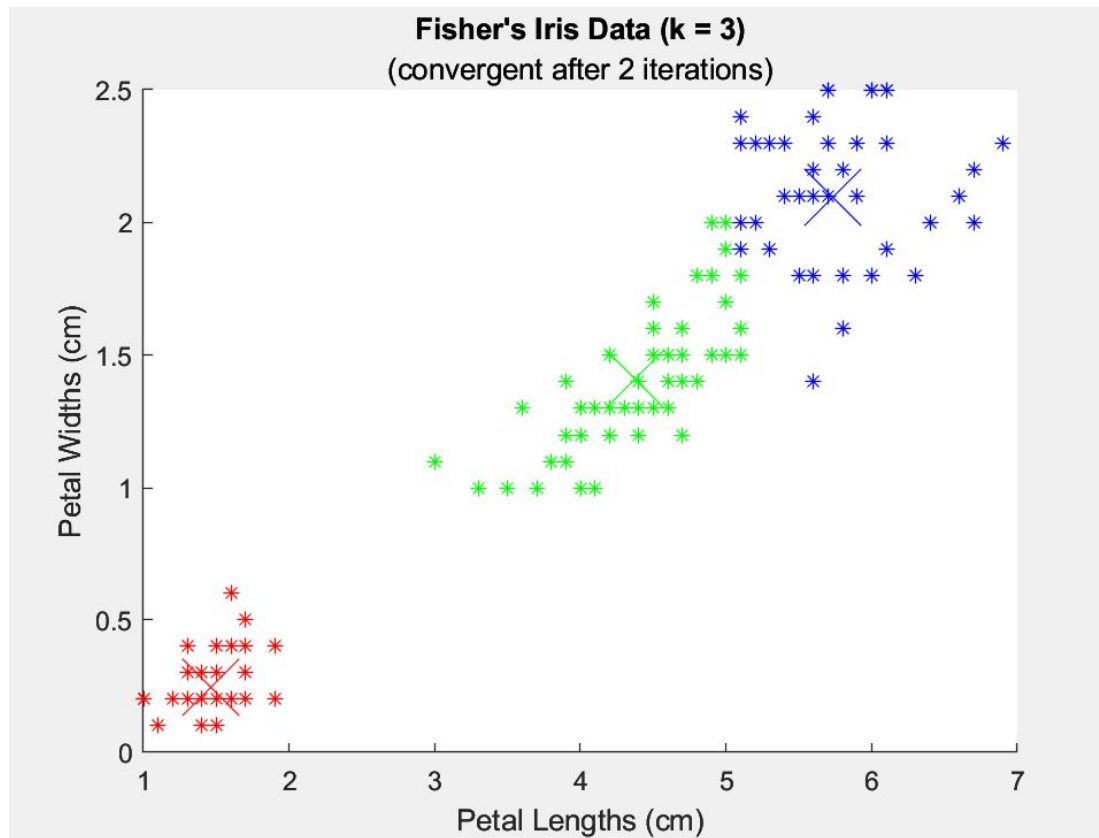
        Assignment step

        Update step

The algorithm *converged* if in the update step the sum of the distances between the old and new cluster center points is less than a threshold:

$$\sum_{j=1}^{k} \left\| \mu_j^{(t+1)} - \mu_j^{(t)} \right\|^2 < \varepsilon \qquad \varepsilon = 0.02$$

Exercise 4 – result

Fisher's Iris Data (k = 3)
(convergent after 2 iterations)

Let's work with image data

An RGB color image is represented by a 3D matrix. It has h rows, w columns and 3 layers along the 3rd dimension. These layers are the R, G and B color layers.

To be able to cluster an image with the the k-means function we *somehow* has to transform the 3D matrix into a 2D array of row vectors.f
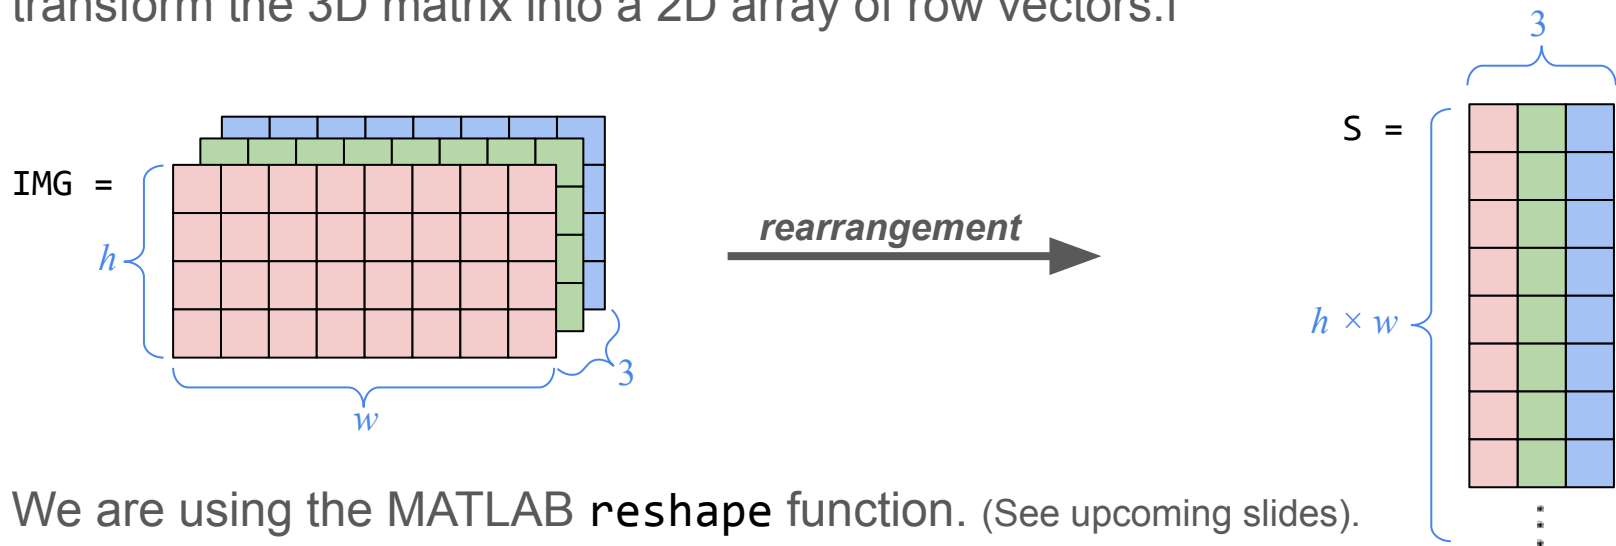

We are using the MATLAB `reshape` function. (See upcoming slides).

How to create a segmented image?

Before clustering

S =

x_1
x_2
x_3
x_4
x_5
x_6
x_7
x_8

Result of the clustering

k = 2

Index: 1   2   3   4   5   6   7   8

LUT =	1	1	1	2	2	2	1	2

M =

mu_1
mu_2

Replacement result

A =

a_1 = mu_1
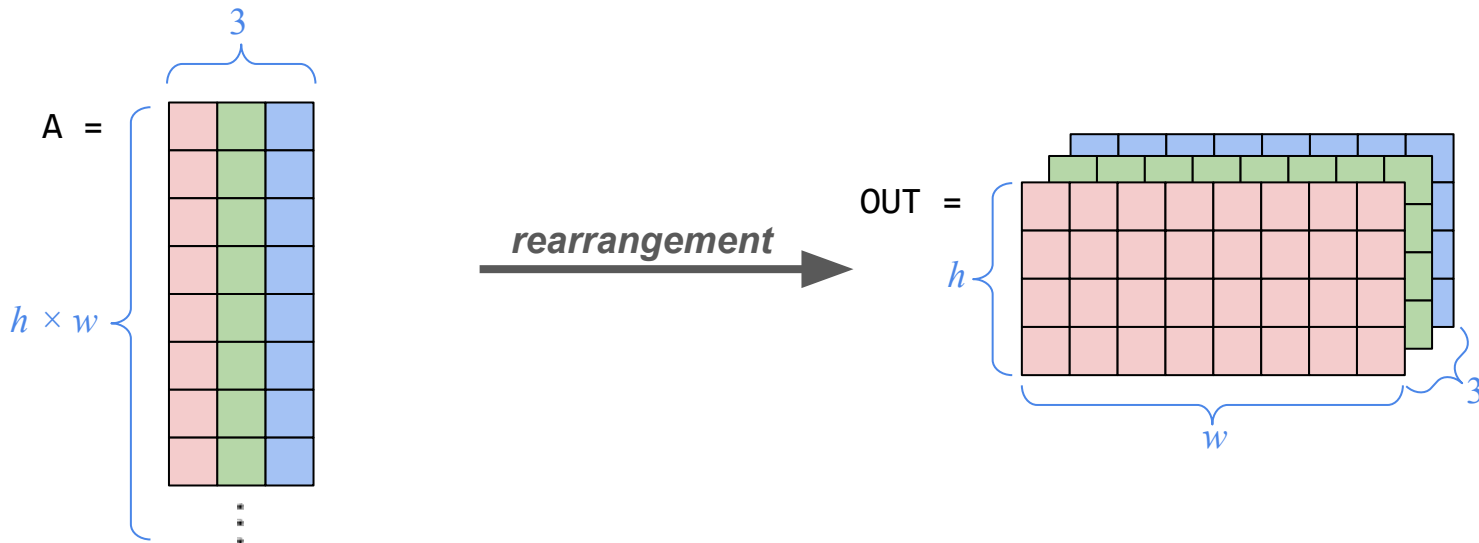a_2 = mu_1
a_3 = mu_1
a_4 = mu_2
a_5 = mu_2
a_6 = mu_2
a_7 = mu_1
a_8 = mu_2

Let's work with image data

After the segmentation is complete we should rearrange the resulted A matrix into a matrix which has the sizes of the original image. This operation is the reverse of the transformation seen on the previous slide.

Exercise 5

Implement the function `image_segmenter` in which:

- The function has 2 inputs and 1 output:

  Inputs:

  - `I`   the RGB image to be clustered (segmented) in double format
  - `k`   number of clusters

  Outputs:

  - `OUT` the **final** segmented image

The function should convert the image to an `S` representation, call the function `mykmeans` on this array, and using the returned `LUT` and `M` create the out image.

After implementation, test your function with the **script** `test5_segmenter`!

Exercise 5 – result

Input image

RGB segmented image (k=2)

THE END